

Robust Reinforcement Learning

Mathias Beguin

2018

1 Formalization of the problem

Reinforcement learning (abbreviated RL) is an area of machine learning in which agents learn how to interact in an environment so as to maximize a notion of reward. It differs from standard supervised learning in that correct input/output pairs need not be presented, and sub-optimal actions need not to be explicitly corrected. Instead it focuses on performance which involve finding a balance between exploration of uncharted state and exploitation of current knowledge.

1.1 Markov Decision Process

In reinforcement learning, the problem is usually described as a Markov Decision Process. It is a mathematical framework used to model decision making problem in which outcomes are under the control of an agent. It is defined as a 5-tuple (S, A, P_a, R_a, γ) where

- S is a set of states
- A is a set of actions (A_s is the set of available actions in state s)
- $Pr(s_{t+1} = s' | s_t = s, a_t = a)$ is the probability of transition from state s to state s' under the action a at time t
- $R(s, a, s')$ is the immediate reward after a transition from s to s' by taking the action a
- $\gamma \in [0, 1)$ is the discount factor, which represents the difference in importance between future and present rewards

A reinforcement learning agent interacts with its environment in discrete time steps. At each time t , the agent receives an observation o_t . It then chooses an action a_t from the set of available actions, which is subsequently sent to the environment. The environment moves to a new state s_{t+1} and the reward r_t associated with the transition (s_t, a_t, s_{t+1}) is determined.




Figure 1: Reinforcement learning interaction loop

The goal of a reinforcement learning agent is to learn a policy π that specifies the action $\pi(s)$ to take in state s . This policy should maximize some cumulative function of the rewards. Typically the expected discounted sum of rewards $\sum_{t=0}^{\infty} \gamma^t R(s, a_t, s')$ where $a_t = \pi(s)$ is taken.

1.2 Policy Gradient

Policy based models are methods that learn a parameterized policy $\pi(a|s;\theta)$ that can select actions without consulting a value function. A value function may still be used to learn the policy parameter, but is not required for action selection. They work by computing an estimator

$$\hat{g} = \hat{\mathbb{E}}_t[\nabla_{\theta} \log \pi_{\theta}(a_t|s_t)\hat{A}_t] \quad (1)$$

using trajectories generated by the policy and plugging it into a stochastic gradient ascent algorithm. The term \hat{A}_t corresponds to an estimator of the advantage function at time step t . It has for purpose to determine which action yields better or worse result than the average. When an estimator of the value function $V_{\pi}(s) = E[\sum_{t=0}^{\infty} \gamma^t r_t | s_0 = s]$ is used, the methods become *actor-critic methods*.

1.3 Contribution

The goal of this thesis is to augment the initial goal of learning a parametrized policy so as to make it robust to the value of nuisance parameters $z \in Z$. Inspired by the work of Gilles Louppe, Michael Kagan and Kyle Cranmer in "Learning to Pivot with Adversarial Networks", our data generation process is a Markov Decision Process defined as in section 1.1 where we modify the transition probability distribution to introduce the nuisance parameters and it becomes $Pr(s_{t+1}|s_t, a_t, z)$.

To learn such policy, we need it to be a pivot with respect to Z . A pivot is a quantity whose distribution is invariant with the nuisance parameters. Therefore, we would like to learn policies such that

$$\pi(a|s, z; \theta) = \pi(a|s, z'; \theta) \quad (2)$$

for all $z, z' \in Z$, all values of $a \in A$ and all values of $s \in S$. This equation implies that $\pi(a|s, \theta)$ and Z are independent random variables.

1.4 Method

Following the method of "Learning to Pivot", we use a re-purposed adversarial networks as a means to constrain the policy π . We pit π against an adversarial model r that takes as input a sequence of tuples $\tau = (s_t, a_t, r_t, s_{t+1})$ where $t \in [0, T]$ which correspond to an episode and where the a_t are chosen using $\pi(a|s; \theta)$ and outputs a function modeling the posterior probability z given τ :

$$r := Pr(z|\tau \sim \pi(a|s; \theta)) \quad (3)$$

Intuitively, if $\pi(a|s, z; \theta)$ varies with z , then the corresponding correlation can be captured by r . By contrast, if $\pi(a|s, z; \theta)$ is invariant with z , then r should perform poorly and be close to random guessing.